

The U-Shaped Challenge of Language Grounding: How Instruction Granularity Induces VLA Shortcuts

Anonymous ACL submission

Abstract

Instruction granularity remains an uncontrolled variable in language-guided embodied AI, leaving its impact unknown. We propose an agent-centric formalization of granularity as rule width (w), a cross-task metric adapted from width-based planning. By measuring the peak concurrent state-feature dependencies an agent must resolve, we introduce w as a measure of instruction granularity and show that it consistently correlates with the decision-making burden neural agents face when grounding language. We also introduce MINI-BEHAVIOR-GRAN, a controlled benchmark systematically varying w across 20 tasks while holding perception and control constant. Our analysis reveals a non-monotonic, U-shaped performance curve: while fine-grained instructions ($w \leq 1$) enforce deep semantic alignment, coarse ones ($w \geq 2$) induce shallow grounding and visuo-motor shortcuts driven by high instruction predictability ($P(l|v)$). We further uncover a generalization asymmetry, where coarse-trained models transfer to fine instructions but not vice-versa, and show that mixed-granularity training is essential for robust language grounding.

1 Introduction

Despite the rapid progress in language-guided embodied AI, the community has yet to establish a consensus on a unified metric to quantify the granularity of instructions during language grounding. In this domain, language serves a specific functional role: it provides structural information to decompose complex tasks and reduce the search space of the underlying sequential decision-making process. Consequently, *instruction granularity* should not be conflated with superficial syntactic properties, such as sentence length or the number of entities mentioned; rather, it must rigorously reflect the *decision-making burden* that remains for the agent to resolve internally after grounding the language into the action space.

Existing research conceptualizes instruction granularity through two dominant lenses, both of which struggle to provide a consistent measure of the decision-making burden required for grounding. Studies in Natural Language Processing (NLP) typically operationalize granularity via surface-form properties, such as textual specificity or discourse structure (e.g., number of entities mentioned) (Dong and Lapata, 2018; Huang et al., 2025b). While these effectively measure syntactic complexity, they adopt a *text-centric* rather than *agent-centric* view. Consequently, they cannot quantify how much an instruction reduces an embodied agent’s *decision-making burden*.

From an *agent-centric* perspective, the essential function of instruction is to decompose complex tasks into tractable sub-problems, thereby reducing the planning burden (Drexler et al., 2022). Recent work in LLM-based agents thus treats granularity as *decomposition depth*—the extent to which an instruction enumerates subgoals versus stating only the final objective (Yang et al., 2024; Myers et al., 2024; Shi et al., 2025; Zhong et al., 2025). While this aligns with our view that *instruction granularity* modulates planning demand, decomposition depth remains an unstable, task-dependent proxy. It conflates instructional guidance with the task’s inherent structure: complex tasks with interleaved subgoals may resist breakdown (e.g., the Sussman Anomaly (Sussman, 1975)), while trivial ones can be artificially elongated.

To operationalize this, we draw inspiration from the concept of policy sketches (Drexler et al., 2024) and formalize instruction granularity as rule width (w) (Bonet et al., 2019). We hypothesize that w captures the decision-making burden that an instruction imposes, and we provide initial empirical evidence supporting this hypothesis, showing that it is a more reliable indicator than common linguistic heuristics. To this end, the contributions of our paper are:

(1) Width as a Formalism for Instruction Granularity: We formalize instruction granularity as *rule width* (w), adapting a concept from width-based planning (Bonet et al., 2019; Drexler et al., 2024) that measures the peak number of concurrent state variables (features Φ) that must be tracked to fulfill the instructions. We argue w measures the *intrinsic decision complexity* of a task-instruction pair, applicable to both symbolic and neural agents (§ 3).

(2) A Controlled Benchmark: We present MINI-BEHAVIOR-GRAN, the first testbed to isolate instruction granularity from confounding factors like task horizon and perception. It enables the first study of how instruction granularity affects VLA performance (§ 4).

(3) The U-Shaped Trend and Width Cap: In this controlled benchmark, we identify a non-monotonic U-shaped relationship where performance peaks at the Fine ($w = 0$) and Coarse ($w \geq 2$) extremes. We also identify a performance collapse at high width ($w \geq 4$): agents fail to fulfill instructions requiring maintaining concurrent features ≥ 4 , revealing a sharp language grounding boundary within our experimental setting (§ 5).

(4) Shallow Grounding and Visuo-Motor Shortcuts: We diagnose the performance rebound at $w \geq 2$ not as true mastery, but as a shift toward *shallow grounding*. Driven by high instruction predictability ($P(l|v)$), agents increasingly rely on visuo-motor shortcuts rather than deep semantic alignment. Through controlled semantic perturbations, we show that mixed-granularity training mitigates this vulnerability, equipping agents with deeper semantic resilience against suboptimal instructions. Additionally, we identify an *asymmetric generalization* pattern: agents trained on coarse instructions successfully transfer to fine-grained ones, but not vice versa.

2 Related Work

Instruction Abstraction and Granularity. In this work, *instruction granularity* refers to the extent to which state and action constraints are explicitly specified versus left for the agent to infer when grounding language into actions. This is distinct from, yet often conflated with the concept of *linguistic abstraction*. The term *abstraction* is viewed as *representational compression*—using a compact description to stand for multiple instances or executions. This definition is common in prior work of

computational thinking (Wing, 2010; Lachmy et al., 2022; Zheng et al., 2024). Importantly, granularity can vary independently of abstraction: “tidy the room” is coarse but non-abstract (it leaves the exact required actions underspecified, yet it lacks representational compression.), whereas “repeat wiping until clean” is relatively specific yet abstract (it prescribes an action but uses a loop condition to compress multiple executions).

While not framed as the term “granularity”, we acknowledge that a rich body of work has focused on augmenting language specificity and its impact, for example through coarse-to-fine semantic parsing (Dong and Lapata, 2018; Li et al., 2020). Other work generates compositional instructions via constraint dropout and sub-task recombination (Huang et al., 2025b), or construct counterfactual instructions to augment data diversity (Glossop et al., 2025).

In embodied instruction following, varying instruction detail is often handled via a simple two-part decomposition: high-level plans are converted into sequences of low-level steps before being grounded into actions (Yang et al., 2024; Shi et al., 2025). Most recent VLA works, however, operate with static language instructions (Ma et al., 2025; Li et al., 2025). Some models like InstructVLA and ThinkAct explore having chain-of-thought reasoning steps before action decoding (Yang et al., 2025; Huang et al., 2025a). These hierarchical definitions remain task-specific and qualitative, lacking a unified, cross-task metric to establish shared granularity levels across tasks. Consequently, quantitatively controlled studies on the formal dynamics of instruction granularity remain scarce in the embodied AI domain.

Language Grounding vs. Shortcuts. Prior work highlights a dichotomy between fragile language sensitivity (Akula et al., 2022) and visuo-motor shortcuts (Xing et al., 2025). We reconcile this via rule width (w): low w exhibits low predictability ($P(l|v)$), enforcing deep language grounding. Conversely, $w \geq 2$ yields high $P(l|v)$: this redundancy induces a shift toward shallow grounding, where agents default to visual priors.

3 Rule Width as Instruction Granularity

3.1 Running Example

We first illustrate *planning width* computation using a running example. Table 1 shows the desired state progressions of a task where two shoes are

State Progression	Concurrent Features Tracked (Sequence of Tuples)	Width
Initial State	$[\langle \rangle]$	0
Navigate to rag	$[\langle p_r = X \rangle, \dots, \langle p_r = 0 \rangle]$	1
Pick up rag	$[\langle p_r = 0 \rangle, \langle H_r \rangle]$	1
Navigate to sink	$[\langle H_r, p_s = X \rangle, \dots, \langle H_r, p_s = 0 \rangle]$	2
Turn on tap (sink)	$[\langle H_r, p_s = 0 \rangle, \langle T_s \rangle]$	2
Soak rag	$[\langle H_r, T_s, p_s = 0 \rangle, \langle S_r \rangle]$	3
Navigate to Shoe 1	$[\langle H_r, S_r, p_{f1} \downarrow \rangle, \dots]$	3
Clean Shoe 1	$[\langle H_r, S_r, p_{f1} = 0 \rangle, \langle C_1 \rangle]$	3
Navigate to Shoe 2	$[\langle H_r, S_r, C_1, p_{f2} \downarrow \rangle, \dots]$	4
Clean Shoe 2	$[\langle H_r, S_r, C_1, p_{f2} = 0 \rangle, \langle C_1, C_2 \rangle]$	4
Put rag on floor	$[\langle C_1, C_2, H_r \rangle, \dots, \langle C_1, C_2, \text{onfloor}_r \rangle]$	3

Table 1: State progression and width for the “Cleaning k Shoes” ($k = 2$). Width scales as $w = k + 2$.

cleaned using a rag, and the rag is placed on the floor afterwards. The key **environment dynamics** are: (1) the rag must be soaked before it can clean shoes, and (2) the tap (sink) must be turned on to soak the rag. We represent states using a set of task-relevant feature variables (Φ): (i) boolean features H_r (holding rag), T_s (tap (sink) is turned on), S_r (rag is soaked), C_i (shoe i is cleaned), onfloor_r (rag is on the floor); and (ii) spatial features p_r, p_s, p_f (distance to the rag, sink, and nearest uncleaned shoe, respectively).

First, approaching the rag is a navigation-only segment: the desired state progression is to decrease the distance feature p_r until $p_r = 0$. Thus, the width is 1 as only a single feature needs to be tracked. In this environment dynamics, $p_r = 0$ enables the state transition to H_r . Thus, the rag can be obtained at $p_r = 0$ and the width remains 1. Subsequently, H_r must be maintained while progressing to next states, resulting in a width of 2. The rag becoming soaked requires simultaneous maintenance of three features, i.e., at the sink ($p_s = 0$), holding the rag (H_r), and having the tap (sink) turned on (T_s), which yields a width of 3. After soaking, C_1 can be obtained by tracking H_r, S_r , and p_f , giving width 3. The subtlety arises when cleaning the second shoe: one has to track the same features as before, plus maintaining the cleanliness of the first shoe (C_1), which increases the width to 4. A special case of $w = 0$ occurs for situations where the desired state progression can be achieved in zero or one step (Bonet and Geffner, 2024); this will appear in our experimental analysis.

Why no action model? Agents’ action sets constrain how state transitions are executed, but *width* focuses on state-level feature dependencies induced by the environment dynamics. Different agents (e.g., a rover vs. a drone) may have different action spaces to achieve the same feature change, so

emphasizing state progressions promotes generalization of width-based concepts across different embodiments.

The impact of *instruction granularity* on VLA behavior remains underexplored largely because the field lacks a **formal, quantifiable, cross-task** measure of granularity. Building on the limitations of heuristic metrics discussed in § 1, we formally equate *instruction granularity* with *latent planning demand*, i.e., the implicit decision-making burden required to ground text into executable actions. We operationalize this using *rule width* (w) by framing natural-language instructions as policy sketches (Drexler et al., 2024). Under this framework, fine-grained instructions decompose tasks into minimal-width sub-problems, whereas coarse instructions leave a large “grounding gap” that the agent must resolve internally.

3.2 Symbolic Task Representation

While VLAs operate on raw multimodal observations, we adopt a symbolic, relational state abstraction. Many embodied benchmarks provide such structured representations—object properties and relations—that can be expressed as *fluents* to form symbolic state or goal conditions (Shridhar et al., 2020; Li et al., 2024). Formally, we consider a classical planning problem $P = \langle \mathcal{F}, \mathcal{A}, I, G \rangle$, where $\mathcal{F}$ is a set of fluents, $\mathcal{A}$ a set of actions, I the initial state, and G the goal.

To bridge raw states and high-level instructions, we define a set of **features** $\Phi = \{\phi_1, \dots, \phi_n\}$, where ϕ_i is a Boolean or integer-valued function over state space S . These features serve as domain-level abstractions that capture task-relevant aspects of the environment. Thus, each s is associated with a **feature valuation vector** $\phi(s) = (\phi_1(s), \dots, \phi_n(s))$.

3.3 Instructions as Rule Sets $\mathcal{R}$

We formalize each instruction as a symbolic interpretation, represented as a set of “if-then” rules $\mathcal{R} = \{r_1, \dots, r_m\}$ defined over the feature space Φ . This rule-based representation does not aim to capture the full expressiveness of natural language; instead, it distills the constraint/goal structure that governs the computational cost incurred by a symbolic search algorithm (i.e., planning demand). Specifically, each r_i is of the form $C_{r_i} \mapsto E_{r_i}$:

- C_{r_i} (the *condition*) is a conjunction of feature constraints. For a Boolean feature p , a constraint




Figure 1: (Left) Factoring r_1 into r_{1a} - r_{1c} reduces latent planning complexity via explicit action guidance. (Right) Fine instructions segment the long-horizon coarse trajectory (π_{r_1} , red).

is p or $\neg p$; for a numerical feature n , it is $n = 0$ or $n > 0$.

- E_{r_i} (the *effect*) is a conjunction of feature changes. For a Boolean feature p , an effect is p , $\neg p$, or $p?$ (allowing any change); for a numerical feature n , it is $n \downarrow$ (decrease), $n \uparrow$ (increase), or $n?$ (any change).

These forms capture the core semantics of the instructions we study. A rule r_i is **fulfilled** by a state sequence $(s, \dots, s')$ if $\phi(s) \models C_{r_i}$ and the pair $(\phi(s), \phi(s'))$ satisfies the changes in E_{r_i} (e.g., $\forall(n \downarrow) \in E_{r_i}, \phi_n(s') < \phi_n(s)$) while keeping unmentioned features constant (i.e., $\phi_j(s') = \phi_j(s)$ for $j \notin E_{r_i}$). To ensure that the agent’s performance strictly reflects its grounding ability rather than the unsolvability of the task, we guarantee that all instruction rule sets $\mathcal{R}$ are executable and dead-end-free. Formally, this is achieved by ensuring the rules are well-formed (i.e., they terminate at the goal) and satisfy the soundness and closure properties (Bonet and Geffner, 2021).

3.4 Rule Width as a Granularity Measure

The width $w(r_i)$ of a rule r_i represents the number of state variables (features) that must be tracked concurrently to fulfill the transition. Formally, for a rule to have width k , the agent must coordinate at most k features from Φ to reach the effect E_{r_i} from a state satisfying C_{r_i} (Lipovetzky and Geffner, 2012).

We define the granularity of an instruction $\mathcal{R}$ by its width-based planning bottleneck: the maximum width among its constituent rules: $w(\mathcal{R}) = \max_{r_i \in \mathcal{R}} w(r_i)$.

We continue the running example to demonstrate two instruction sets, each representing a different width level, with two variants visualized in Figure 1.

We also append a **counting feature** in Φ to track the number of uncleaned shoes, denoted as N_f .

Coarse-Grained Instruction ($w(\mathcal{R}) = 3$) This instruction states only the high-level rules and their desired ordering, leaving the internal structure of each sub-task implicit:

- $r_1 : \{N_f > 0\} \mapsto \{N_f \downarrow\}$ (clean a shoe; $w = 3$)
- $r_2 : \{N_f = 0, \neg on_floor_r\} \mapsto \{on_floor_r\}$ (put the rag on the floor after cleaning shoes; $w = 1$)

Middle-Grained Instruction ($w(\mathcal{R}) = 2$) This variant decomposes r_1 to explicitly state the *soaking* process:

- $r_{1a} : \{\neg H_r\} \mapsto \{H_r\}$ (pick up the rag; $w = 1$)
- $r_{1b} : \{H_r, \neg S_r\} \mapsto \{S_r, H_r\}$ (soak the rag)
- $r_{1c} : \{N_f > 0, S_r\} \mapsto \{N_f \downarrow\}$ (clean; $w = 2$)
- $r_2 : \{N_f = 0, \neg on_floor_r\} \mapsto \{on_floor_r\}$ (put the rag on the floor; $w = 1$)

This rule set is well-formed: it is goal-terminating via r_2 , and the rules are sequentially reachable, meaning the successor state of any fulfilled rule (that is not the goal) also meets the condition of at least one rule in $\mathcal{R}$. The set is also robust to perturbations: if an agent drops the rag after soaking it ($\neg H_r$), the rule set remains applicable by reverting to r_{1a} to reacquire the rag before proceeding to the cleaning task (r_{1c}).

Width as a Planning Proxy. Decomposing a task into a rule set $\mathcal{R}$ provides structural scaffolding, guiding the agent through intermediate “stepping stone” states. By design, we assume the execution of r_i does not negate established fluents unless explicitly required by E_{r_i} . This is the essence of language-guided decomposition. Without this assumption, language guidance would offer no reduction in complexity; the agent would need to track every feature since the initial state.

This intuition is formally grounded in novelty-driven Breadth-First Search (Lipovetzky and Geffner, 2017). By pruning states that fail to achieve a new feature conjunction of size $\leq k$, search complexity is bounded by $O(|\Phi|^k)$. Thus, trajectories that unnecessarily negate an already-achieved precondition C_{r_i} (unless required by E_{r_i}) are longer than those preserving it and are pruned from the optimal frontier. Consequently, features in C_{r_i} that are already true are effectively treated




Figure 2: MINI-BEHAVIOR-GRAN Experimental Pipeline. We evaluate the impact of instruction granularity across three VLA action-decoding strategies. During inference time, an Oracle Instructor evaluates the current world state to decide available rule set $\mathcal{R}_a$ and active rule r_t . A rule-based natural language instruction l_t is generated accordingly. The process iterates until completion or max steps.

as invariants, and the active coordination w reduces to the incremental features needed for the next sub-goal. For instance, while the monolithic task requires coordinating $\{H_r, p_s, T_s\}$ ($w = 3$, Table 2), the instruction r_{1b} leverages the fact that H_r is pre-established. By treating H_r as a fixed invariant, the search width w reduces to the novel features $\{p_s, T_s\}$, yielding $w = 2$.

To communicate these rules to VLA agents, we translate each rule into a natural-language sentence via a template-based generator. These templates verbalize the features within each rule; e.g., $\{H_r, \neg S_r\} \mapsto \{S_r\}$ is realized as “If you are holding the rag but the rag is not yet soaked, soak the rag.” This construction preserves the granularity differences defined by the rule width. Additional language prompt examples are provided in § F.

4 Mini-BEHAVIOR-Gran Benchmark

To disentangle the influence of instruction granularity in VLA performance, we seek an environment that (1) retains long-horizon and heterogeneous activities so that instruction guidance meaningfully modulates the planning demand, while (2) simplifies perception and low-level control to avoid confounding factors. As such, we use MINI-BEHAVIOR (Jin et al., 2023) as our testbed. Detailed justification for our choice over other environments is provided in § B.

We introduce MINI-BEHAVIOR-GRAN, the first benchmark enabling controlled study of instruction granularity via *rule width*. Over 20 tasks sharing a feature set Φ , we design rule sets at up to

six width levels, all verified solvable. During inference, an oracle instructor monitors the current feature vector $\phi(s_t)$, selects an available rule r_t from the task’s rule set, and translates it via a template into a natural language instruction l_t . This guarantees that performance differences reflect grounding ability alone, not stochastic instruction quality (see Figure 2). The instruction l_t and current observation v_t are then fed to the VLA model to predict action chunks $(a_t, \dots, a_{t+k}) \sim \pi_{\text{vla}}(\cdot | v_t, l_t)$.

5 Experiments

5.1 Experimental Setup

Task Suite & Training. We use MINI-BEHAVIOR-GRAN, which contains 20 complex tasks sharing the same feature set Φ . We further classify instruction granularity into three *relative* classes of *Fine* (F), *Medium* (M), and *Coarse* (C) via per-task quantile binning. In most cases, this mapping corresponds to $w = 0$ for Fine, $w = 1$ for Medium, and $w \geq 2$ for Coarse. We evaluate seven training settings: (1) **Pure F/M/C** ($1 : 0 : 0$): trains exclusively on one granularity class; (2) **Centroid** ($1 : 1 : 1$): uniformly sample from all three classes; (3) **Axial F/M/C** ($3 : 1 : 1$): emphasizes one class with some exposure to the other classes. Our dataset provides 50 training and 10 evaluation instances per task, each paired with reference trajectories generated by BFWS width-based planner from Lipovetzky and Geffner (2017) (see § E for details).

Action-Decoding Archetypes. The primary architectural difference among latest VLA models




Figure 3: Rule call distribution across instruction granularities. # of rule calls decreases with coarser instructions. It decouples instruction length from width.




Figure 4: Impact of width and horizon on SR. SR decays with task horizon; larger instruction width imposes an additional burden even when horizon is fixed.

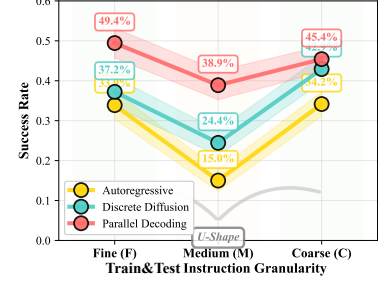


Figure 5: In-Distribution (matched train/test) SR across granularities (i.e., train Fine $\rightarrow$ test Fine, etc.). A non-monotonic U-shaped trend appears.

lies in their action-decoding strategies (Zhong et al., 2025). Hence, we employ a unified VLM backbone but vary three representative action-decoding strategies (see Figure 2 left panel): (1) **Autoregressive** (AR), exemplified by RT-2 (Zitkovich et al., 2023), OpenVLA (Kim et al., 2024) and others (Reed et al., 2022; Huang et al., 2024; O’Neill et al., 2024) that predict actions sequentially via causal attention; (2) **(Discrete) Diffusion**, which uses diffusion process for action chunk prediction (Black et al., 2025; Bjorck et al., 2025; Shukor et al., 2025). We employ discrete diffusion (Liang et al., 2025) for compatibility with the discrete action space of our testbed; (3) **Parallel Decoding**, utilizing bidirectional attention for single-pass action prediction to improve throughput, as demonstrated by OpenVLA-OFT (Kim et al., 2025b), PD-VLA (Song et al., 2025) and other variants (Yin et al., 2025; Wang et al., 2025; Xiao et al., 2025). Details on model hyperparameters, hardware and training schedules are in § G.

5.2 Establishing the Yardstick: Decoupling Horizon from Width

A critical confounder in evaluating VLA planning complexity is the task sequence length. However, in language-guided execution, instructions inherently decompose long-horizon tasks into segments. Consequently, the relevant measure of sequential burden shifts from the raw atomic horizon (total action steps) to the instruction horizon, i.e., the number of rule calls issued by the instructor.

To isolate rule width from sequential burden, we train a single agent across all tasks using the Centroid configuration, ensuring balanced exposure across all granularity levels. This approach holds the distribution of atomic horizons constant, allow-

ing us to isolate how instruction granularity and the frequency of rule calls affect performance. Figure 3 shows that fine-grained instructions ($w = 0$) decompose the task into numerous low-width rule calls; coarse instructions ($w \geq 2$), by contrast, yield fewer rule calls.

To assess whether w serves as an effective indicator of latent planning demand for learning-based agents, we analyze Figure 4. The heatmap reveals a strict monotonic decay in Success Rate (SR) as rule width increases. Also, for a fixed w , SR decays as rule calls increase, mirroring greater sequential horizon complexity. This confirms w as a horizon-independent measure of the intrinsic planning demand imposed by instruction granularity.

Width vs. Syntactic Heuristics. We evaluate whether w captures language grounding complexity better than common linguistic heuristics: token count, entity count, and action verbs. Figure 8 shows w is the only metric that exhibits a statistically significant and consistently *negative* correlation with SR across task horizons. Action Verbs correlate significantly but flip direction across horizons, while Token Count turns negative only in the longest bin (20+ steps). These heuristics are inconsistent compared to w (see details in § H).

5.3 The U-Shaped Paradox and Capacity Cap

To establish a baseline, we train models using the *Pure F/M/C* setting. the models are trained and evaluated on the same granularity class. As shown in Figure 5, all VLA archetypes exhibit a non-monotonic *U-shaped performance curve*.

At the fine extreme ($w = 0$), where instructions require minimal latent reasoning to be grounded into actions, models achieve high performance. Performance then decreases at medium granularity




Figure 6: Smaller SR gaps indicate a shift towards shallow grounding. Under semantic perturbation, the Pure-C setting exhibits tiny variance (2.8%) between w/ goal lang and w/o lang.




Figure 7: VLM instruction-prediction Cross Entropy (CE) loss across granularities and 3 seeds. Larger CE loss indicates lower $P(l | v)$.




Figure 8: Metric-success correlation by task-horizon.

($w = 1$), before unexpectedly rebounding at the coarse regime ($w \geq 2$). This contradicts the assumption that higher planning width should monotonically degrade learning.

5.4 Mechanistic Probe: The Grounding Valley

Why does performance rebound at $w \geq 2$? We show this arises from increasing *instruction predictability* $P(l|v)$, which enables agents to bypass language and exploit visual-dominant policy.

Following prior work (Fei et al., 2025), we use $\Delta SR_{w/-w/o}$ as a basic measure of language grounding. To further isolate grounding effects, we conduct another intervention: replacing full instructions with high-level static goal descriptions (Figure 6). When language is entirely removed, fine-trained policies suffer a severe 36.3% drop, whereas coarse-trained policies show a smaller 20.0% gap. Also note that under reduced informativeness (w/ goal lang), the Pure-C policy exhibits a negligible 2.8% advantage over having no language at all (15.6% vs. 12.8%). This indicates that policies trained on coarse instructions develop a shallow grounding behavior: language is not deeply integrated into their decision-making. Con-

sequently, when instructions are reduced to generic goals, their performance barely differs from the vision-only baseline.

The Predictability Driver: Why Reliance Shifts.

We hypothesize that this behavioral shift originates from coarse instructions being statistically more predictable from vision. Formally, consider a surjective mapping f from fine-grained labels L_{fine} to coarser ones L_{coarse} . The conditional probability of a coarse label aggregates its fine-grained constituents: $P(c_j|v) = \sum_{l_i \in f^{-1}(c_j)} P(l_i|v)$. The increasing predictability of coarse instructions reduces the marginal information contributed by language, making it easier for agents to minimize task loss by relying primarily on visual observations, thereby favoring the development of visuo-motor shortcuts over costly cross-modal alignment. We confirm this predictability empirically by training a VLM (Qwen3-VL-4B via LoRA) to act as an trained instructor, predicting instructions from visual observations. As Figure 7 shows, Cross-Entropy (CE) loss decreases monotonically with coarseness (Fine: 0.078, Medium: 0.063, Coarse: 0.047), a relative drop of 40% from Fine to Coarse. Together, these results form a complementary chain: probability aggregation explains *why* $P(l|v)$ increases, CE loss empirically confirms the increase, and the perturbation ΔSR quantifies the resulting visuo-motor shift.

Thus, we attribute this U-shaped behavior to how the agent dynamically shifts its reliance on text versus vision across different rule widths:

- **Fine** ($w = 0$): *Deep Grounding*. Low $P(l|v)$ (high CE loss, Figure 7) means language provides unique information. Combined with low-width grounding burden (§ 5.2), the model reliably maps instructions to actions, yielding high

Training Setup	Width 0	Width 1	Width 2	Width 3	Width ≥ 4
Pure-F	34.9%	16.9%	10.8%	1.0%	0.0%
Pure-M	12.2%	22.0%	11.7%	0.0%	0.0%
Pure-C	26.3%	33.3%	32.9%	22.9%	0.0%
Axial-F	50.2%	53.3%	48.8%	34.3%	0.0%
Axial-M	44.7%	43.5%	45.8%	29.5%	0.0%
Axial-C	55.7%	52.2%	50.8%	37.1%	0.0%
Centroid (Mix)	42.7%	38.4%	40.8%	28.6%	0.0%

Table 2: The x-axis are test instruction width, y-axis are training setup. Transfer is asymmetric: train coarse $\rightarrow$ test fine works, but train fine $\rightarrow$ test coarse does not. Mixed-granularity improves overall.




Figure 9: Low-width but long task horizon (e.g., laying woods) tends to cause regression, while high-width stagnates.

SR (Figure 5).

- **Medium** ($w = 1$): *The Alignment Trap*. At this width, grounding complexity increases while $P(l|v)$ remains moderate: not yet high enough to exploit reliable visuo-motor shortcuts. The model must attempt language grounding but struggles with its inherent difficulty. This intermediate regime, caught between hard alignment and insufficient shortcuts, yields the lowest performance.
- **Coarse** ($2 \leq w \leq 3$): *Reduced Language Reliance*. High $P(l|v)$ incentivizes a visual-dominant policy: aligning vision to action is easier than learning cross-modal grounding. This yields high in-distribution scores but at the cost of language grounding capability.
- **Capacity Cap** ($w \geq 4$): *Cognitive Overload*. Mirroring the $O(|\Phi|^4)$ complexity barrier of symbolic planners (Lipovetzky and Geffner, 2017), VLA performance uniformly collapses at $w \geq 4$ (Table 2). This convergence suggests w as a fundamental representational bottleneck shared by both symbolic and neural paradigms.

Deepening Grounding via Axial Training. Having diagnosed this shallow reliance as a byproduct of high predictability, we hypothesize that increasing the diversity of the training language space L should artificially lower the average $P(l|v)$, forcing the agent to structurally re-engage with the text. Our mixed-granularity training confirms this mechanism. All mixed-granularity models exhibit substantially larger Δ SR gaps compared to their *Pure-X* counterparts. Among these, the fine-weighted mixture *Axial-F* emerges as the optimal structural compromise: it achieves an overall success rate, while maintaining a significantly larger Δ SR.

5.5 Other Results

Asymmetric Generalization. We observe an asymmetric transfer pattern (Table 2): coarse-trained models successfully zero-shot transfer to fine-grained instructions, but fine-trained models fail on coarse-grained ones. This asymmetry aligns with our Δ SR analysis: coarse training induces a visual-dominant policy, hence perturbing or refining the instruction yields minimal behavioral change. Conversely, fine training induces tight coupling with language, making the policy brittle when detailed guidance is withdrawn.

Failure Modes. Agent failure distributions bifurcate strictly with width (Figure 9). Low-width, long-horizon tasks predominantly suffer from *regression*: the agent violates previously satisfied fluents and cycles back to prior states, reflecting subgoal serialization failures. High-width tasks ($w \geq 3$), by contrast, exhibit *stagnation*: the agent freezes, unable to resolve concurrent feature dependencies ($O(|\Phi|^w)$) required to initiate a valid state transition. This dichotomy mirrors the width-based complexity landscape (see § H for details).

6 Conclusion

This work introduces rule width (w) as a cross-task formalism for instruction granularity and demonstrates that it correlates with language grounding difficulty in learning-based agents. We observed a non-monotonic relationship between granularity and performance, driven by a trade-off between instruction predictability ($P(l|v)$) and latent planning complexity. We demonstrate that mixed-granularity training balances language sensitivity without sacrificing the robust visual reasoning skills learned from coarse trajectories.

7 Limitations

Discrete Action Spaces. We evaluate our framework in environments with discrete action spaces, a common abstraction in language-conditioned embodied AI (e.g., ALFWorld-based agent (Xiong et al., 2025; Kim et al., 2025a)) and LLM-based agent systems (Qin et al., 2024; Patil et al., 2024), where high-level reasoning is decoupled from low-level execution via skill primitives or API calls. This choice is intentional: it isolates symbolic planning from continuous control noise, allowing us to cleanly attribute observed effects to instruction granularity. However, extending this granularity-aware framework to hybrid discrete-continuous or continuous action spaces remains essential for real-world robotics applications, where low-level motor commands interact with high-level linguistic abstractions.

On the Oracle Instructor Setup. A potential concern is that providing the active rule via an Oracle simplifies the task to single-step grounding. However, we argue that the **latent planning demand** w is an intrinsic property of the instruction-to-action mapping. Even when the “what to do” is provided, the agent must still coordinate w concurrent features to achieve the transition. Our findings that performance collapses at $w \geq 4$ despite Oracle guidance (Table 2) demonstrates that the grounding gap is a fundamental bottleneck in the learning process.

References

Arjun R. Akula, Spandana Gella, Aishwarya Padmakumar, Mahdi Namazifar, Mohit Bansal, Jesse Thomason, and Dilek Hakkani-Tur. 2022. ALFRED-L: investigating the role of language for action learning in interactive visual environments. In *EMNLP*, pages 9369–9378. Association for Computational Linguistics.

Johan Bjorck, Fernando Castañeda, Nikita Cherniadev, Xingye Da, Runyu Ding, Linxi Fan, Yu Fang, Dieter Fox, Fengyuan Hu, Spencer Huang, and 1 others. 2025. Gr00t n1: An open foundation model for generalist humanoid robots. *arXiv preprint arXiv:2503.14734*.

Kevin Black, Noah Brown, James Darpinian, Karan Dhabalia, Danny Driess, Adnan Esmail, Michael Robert Equi, Chelsea Finn, Niccolo Fusai, Manuel Y. Galliker, Dibya Ghosh, Lachy Groom, Karol Hausman, brian ichter, Szymon Jakubczak, Tim Jones, Liyiming Ke, Devin LeBlanc, Sergey

Levine, and 16 others. 2025. $\pi_{0.5}$: a vision-language-action model with open-world generalization. In *Proceedings of The 9th Conference on Robot Learning*, volume 305 of *Proceedings of Machine Learning Research*, pages 17–40. PMLR.

Blai Bonet, Guillem Francès, and Hector Geffner. 2019. Learning features and abstract actions for computing generalized plans. In *AAAI*, pages 2703–2710. AAAI Press.

Blai Bonet and Hector Geffner. 2021. General policies, representations, and planning width. In *AAAI*, pages 11764–11773. AAAI Press.

Blai Bonet and Hector Geffner. 2024. General policies, subgoal structure, and planning width. *Journal of Artificial Intelligence Research*, 80:475–516.

Li Dong and Mirella Lapata. 2018. Coarse-to-fine decoding for neural semantic parsing. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 731–742, Melbourne, Australia. Association for Computational Linguistics.

Dominik Drexler, Jendrik Seipp, and Hector Geffner. 2022. Learning sketches for decomposing planning problems into subproblems of bounded width. In *ICAPS*, pages 62–70. AAAI Press.

Dominik Drexler, Jendrik Seipp, and Hector Geffner. 2024. Expressing and exploiting subgoal structure in classical planning using sketches. *J. Artif. Intell. Res.*, 80.

Senyu Fei, Siyin Wang, Junhao Shi, Zihao Dai, Jikun Cai, Pengfang Qian, Li Ji, Xinzhe He, Shiduo Zhang, Zhaoye Fei, Jinlan Fu, Jingjing Gong, and Xipeng Qiu. 2025. Libero-plus: In-depth robustness analysis of vision-language-action models. *Preprint*, arXiv:2510.13626.

Catherine Glossop, William Chen, Arjun Bhorkar, Dhruv Shah, and Sergey Levine. 2025. Cast: Counterfactual labels improve instruction following in vision-language-action models. *Preprint*, arXiv:2508.13446.

Danijar Hafner. 2021. Benchmarking the spectrum of agent capabilities. *arXiv preprint arXiv:2109.06780*.

Chi-Pin Huang, Yueh-Hua Wu, Min-Hung Chen, Yu-Chiang Frank Wang, and Fu-En Yang. 2025a. Thinkact: Vision-language-action reasoning via reinforced visual latent planning. *Preprint*, arXiv:2507.16815.

Hui Huang, Jiaheng Liu, Yancheng He, Shilong Li, Bing Xu, Conghui Zhu, Muyun Yang, and Tiejun Zhao. 2025b. Musc: Improving complex instruction following with multi-granularity self-contrastive training. In *ACL (1)*, pages 10667–10686. Association for Computational Linguistics.

726	Jiangyong Huang, Silong Yong, Xiaojian Ma, Xiongkun	Matthews, and 9 others. 2022b. BEHAVIOR-1K:	783
727	Linghu, Puhao Li, Yan Wang, Qing Li, Song-Chun	A benchmark for embodied AI with 1, 000 everyday	784
728	Zhu, Baoxiong Jia, and Siyuan Huang. 2024. An	activities and realistic simulation. In <i>CoRL</i> , volume	785
729	embodied generalist agent in 3d world. In <i>ICML</i> .	205 of <i>Proceedings of Machine Learning Research</i> ,	786
		pages 80–93. PMLR.	787
730	Emily Jin, Jiaheng Hu, Zhuoyi Huang, Ruohan Zhang,	Wei Li, Renshan Zhang, Rui Shao, Jie He, and Liqiang	788
731	Jiajun Wu, Li Fei-Fei, and Roberto Martín-Martín.	Nie. 2025. Cogvla: Cognition-aligned vision-	789
732	2023. Mini-BEHAVIOR: A procedurally generated	language-action model via instruction-driven routing	790
733	benchmark for long-horizon decision-making in em-	& sparsification. <i>Advances in neural information</i>	791
734	bodied AI . In <i>NeurIPS 2023 Workshop on General-</i>	<i>processing systems</i> .	792
735	<i>ization in Planning</i> .		
736	Siddharth Karamcheti, Suraj Nair, Ashwin Balakrishna,	Xuanlin Li, Kyle Hsu, Jiayuan Gu, Oier Mees, Karl	793
737	Percy Liang, Thomas Kollar, and Dorsa Sadigh. 2024.	Pertsch, Homer Rich Walke, Chuyuan Fu, Ishikaa	794
738	Prismatic vlms: Investigating the design space of	Lunawat, Isabel Sieh, Sean Kirmani, Sergey Levine,	795
739	visually-conditioned language models . In <i>Forty-</i>	Jiajun Wu, Chelsea Finn, Hao Su, Quan Vuong, and	796
740	<i>first International Conference on Machine Learning,</i>	Ted Xiao. 2024. Evaluating real-world robot manipu-	797
741	<i>ICML 2024, Vienna, Austria, July 21-27, 2024</i> .	lation policies in simulation. In <i>CoRL</i> , volume 270	798
		of <i>Proceedings of Machine Learning Research</i> , pages	799
742	Jeonghye Kim, Sojeong Rhee, Minbeom Kim, Do-	3705–3728. PMLR.	800
743	hyung Kim, Sangmook Lee, Youngchul Sung, and		
744	Kyomin Jung. 2025a. ReflAct: World-grounded deci-	Zechang Li, Yuxuan Lai, Yansong Feng, and Dongyan	801
745	sion making in LLM agents via goal-state reflection .	Zhao. 2020. Domain adaptation for semantic pars-	802
746	In <i>Proceedings of the 2025 Conference on Empiri-</i>	ing . In <i>Proceedings of the Twenty-Ninth Interna-</i>	803
747	<i>cal Methods in Natural Language Processing</i> , pages	<i>tional Joint Conference on Artificial Intelligence, IJ-</i>	804
748	33433–33465, Suzhou, China. Association for Com-	<i>CAI 2020</i> , pages 3723–3729. ijcai.org.	805
749	putational Linguistics.		
750	Moo Jin Kim, Chelsea Finn, and Percy Liang.	Zhixuan Liang, Yizhuo Li, Tianshuo Yang, Chengyue	806
751	2025b. Fine-tuning vision-language-action mod-	Wu, Sitong Mao, Liuaoyi, Xiaokang Yang, Jiang-	807
752	els: Optimizing speed and success. <i>arXiv preprint</i>	miao Pang, Yao Mu, and Ping Luo. 2025. Discrete	808
753	<i>arXiv:2502.19645</i> .	diffusion vla: Bringing discrete diffusion to action	809
		decoding in vision-language-action policies. <i>arXiv</i>	810
754	Moo Jin Kim, Karl Pertsch, Siddharth Karamcheti,	<i>preprint arXiv:2508.20072</i> .	811
755	Ted Xiao, Ashwin Balakrishna, Suraj Nair, Rafael	Nir Lipovetzky and Hector Geffner. 2012. Width and	812
756	Rafailov, Ethan Paul Foster, Pannag R. Sanketi, Quan	serialization of classical planning problems. In <i>ECAI</i> ,	813
757	Vuong, Thomas Kollar, Benjamin Burchfiel, Russ	volume 242 of <i>Frontiers in Artificial Intelligence and</i>	814
758	Tedrake, Dorsa Sadigh, Sergey Levine, Percy Liang,	<i>Applications</i> , pages 540–545. IOS Press.	815
759	and Chelsea Finn. 2024. Openvla: An open-source		
760	vision-language-action model. In <i>CoRL</i> , volume 270	Nir Lipovetzky and Hector Geffner. 2017. Best-first	816
761	of <i>Proceedings of Machine Learning Research</i> , pages	width search: Exploration and exploitation in clas-	817
762	2679–2713. PMLR.	sical planning. In <i>AAAI</i> , pages 3590–3596. AAAI	818
		Press.	819
763	Royi Lachmy, Valentina Pyatkin, Avshalom Manevich,	Yueen Ma, Zixing Song, Yuzheng Zhuang, Jianye	820
764	and Reut Tsarfaty. 2022. Draw me a flower: Process-	Hao, and Irwin King. 2025. A survey on vision-	821
765	ing and grounding abstraction in natural language.	language-action models for embodied ai . <i>Preprint</i> ,	822
766	<i>Trans. Assoc. Comput. Linguistics</i> , 10:1341–1356.	<i>arXiv:2405.14093</i> .	823
767	Chengshu Li, Fei Xia, Roberto Martín-Martín, Michael	Vivek Myers, Chunyuan Zheng, Oier Mees, Kuan Fang,	824
768	Lingelbach, Sanjana Srivastava, Bokui Shen, Kent El-	and Sergey Levine. 2024. Policy adaptation via lan-	825
769	liott Vainio, Cem Gokmen, Gokul Dharan, Tanish	guage optimization: Decomposing tasks for few-shot	826
770	Jain, Andrey Kurenkov, Karen Liu, Hyowon Gweon,	imitation. In <i>CoRL</i> , volume 270 of <i>Proceedings</i>	827
771	Jiajun Wu, Li Fei-Fei, and Silvio Savarese. 2022a.	of <i>Machine Learning Research</i> , pages 1402–1426.	828
772	igibson 2.0: Object-centric simulation for robot learn-	PMLR.	829
773	ing of everyday household tasks . In <i>Proceedings of</i>		
774	<i>the 5th Conference on Robot Learning</i> , volume 164	Abby O’Neill, Abdul Rehman, Abhiram Maddukuri,	830
775	of <i>Proceedings of Machine Learning Research</i> , pages	Abhishek Gupta, Abhishek Padalkar, Abraham	831
776	455–465. PMLR.	Lee, Acorn Pooley, Agrim Gupta, Ajay Mandlekar,	832
777	Chengshu Li, Ruohan Zhang, Josiah Wong, Cem Gok-	Ajinkya Jain, Albert Tung, Alex Bewley, Alexander	833
778	men, Sanjana Srivastava, Roberto Martín-Martín,	Herzog, Alex Irpan, Alexander Khazatsky, Anant Rai,	834
779	Chen Wang, Gabrael Levine, Michael Lingelbach,	Anchit Gupta, Andrew E. Wang, Anikait Singh, and	835
780	Jiankai Sun, Mona Anvari, Minjune Hwang, Man-	260 others. 2024. Open x-embodiment: Robotic	836
781	asi Sharma, Arman Aydin, Dhruva Bansal, Samuel	learning datasets and RT-X models : Open x-	837
782	Hunter, Kyu-Young Kim, Alan Lou, Caleb R.	embodiment collaboration. In <i>ICRA</i> , pages 6892–	838
		6903. IEEE.	839

840	Maxime Oquab, Timothée Darcet, Théo Moutakanni,	models for robotics manipulation. <i>arXiv preprint</i>	897
841	Huy V. Vo, Marc Szafraniec, Vasil Khalidov,	<i>arXiv:2506.07530</i> .	898
842	Pierre Fernandez, Daniel Haziza, Francisco Massa,		
843	Alaaeldin El-Nouby, Mido Assran, Nicolas Ballas,	Jeannette M. Wing. 2010. <i>Computational thinking:</i>	899
844	Wojciech Galuba, Russell Howes, Po-Yao Huang,	<i>What and why? The Link</i> .	900
845	Shang-Wen Li, Ishan Misra, Michael Rabbat, Vasu		
846	Sharma, and 7 others. 2024. Dinov2: Learning robust	Lei Xiao, Jifeng Li, Juntao Gao, Feiyang Ye, Yan Jin,	901
847	visual features without supervision. <i>Trans. Mach.</i>	Jingjing Qian, Jing Zhang, Yong Wu, and Xiaoyuan	902
848	<i>Learn. Res.</i> , 2024.	Yu. 2025. Ava-vla: Improving vision-language-	903
		action models with active visual attention. <i>arXiv</i>	904
849	Shishir G. Patil, Tianjun Zhang, Xin Wang, and	<i>preprint arXiv:2511.18960</i> .	905
850	Joseph E. Gonzalez. 2024. Gorilla: Large language		
851	model connected with massive apis. In <i>NeurIPS</i> .	Youguang Xing, Xu Luo, Junlin Xie, Lianli Gao,	906
		Heng Tao Shen, and Jingkuan Song. 2025. <i>Short-</i>	907
852	Yujia Qin, Shihao Liang, Yining Ye, Kunlun Zhu, Lan	<i>cut learning in generalist robot policies: The role of</i>	908
853	Yan, Yaxi Lu, Yankai Lin, Xin Cong, Xiangru Tang,	<i>dataset diversity and fragmentation</i> . In <i>9th Annual</i>	909
854	Bill Qian, Sihan Zhao, Lauren Hong, Runchu Tian,	<i>Conference on Robot Learning</i> .	910
855	Ruobing Xie, Jie Zhou, Mark Gerstein, Dahai Li,		
856	Zhiyuan Liu, and Maosong Sun. 2024. Toolllm: Fa-	Weimin Xiong, Yifan Song, Qingxiu Dong, Bingchan	911
857	cilitating large language models to master 16000+	Zhao, Feifan Song, XWang, and Sujian Li. 2025.	912
858	real-world apis. In <i>ICLR</i> . OpenReview.net.	<i>MPO: Boosting LLM agents with meta plan opti-</i>	913
		<i>mization</i> . In <i>Findings of the Association for Com-</i>	914
859	Scott E. Reed, Konrad Zolna, Emilio Parisotto,	<i>putational Linguistics: EMNLP 2025</i> , pages 3914–	915
860	Sergio Gómez Colmenarejo, Alexander Novikov,	3935, Suzhou, China. Association for Computational	916
861	Gabriel Barth-Maron, Mai Gimenez, Yury Sulsky,	Linguistics.	917
862	Jackie Kay, Jost Tobias Springenberg, Tom Eccles,		
863	Jake Bruce, Ali Razavi, Ashley Edwards, Nicolas	Ruochu Yang, Fumin Zhang, and Mengxue Hou. 2024.	918
864	Heess, Yutian Chen, Raia Hadsell, Oriol Vinyals,	Oceanplan: Hierarchical planning and replanning	919
865	Mahyar Bordbar, and Nando de Freitas. 2022. A	for natural language AUV piloting in large-scale un-	920
866	generalist agent. <i>Trans. Mach. Learn. Res.</i> , 2022.	explored ocean environments. In <i>WUWNet</i> , pages	921
		11:1–11:5. ACM.	922
867	Lucy Xiaoyang Shi, Brian Ichter, Michael Robert Equi,		
868	Liyiming Ke, Karl Pertsch, Quan Vuong, James Tan-	Shuai Yang, Hao Li, Yilun Chen, Bin Wang, Yang Tian,	923
869	ner, Anna Walling, Haohuan Wang, Niccolo Fusai,	Tai Wang, Hanqing Wang, Feng Zhao, Yiyi Liao, and	924
870	Adrian Li-Bell, Danny Driess, Lachy Groom, Sergey	Jiangmiao Pang. 2025. <i>Instructvla: Vision-language-</i>	925
871	Levine, and Chelsea Finn. 2025. Hi robot: Open-	<i>action instruction tuning from understanding to ma-</i>	926
872	ended instruction following with hierarchical vision-	<i>nipulation</i> . <i>Preprint</i> , arXiv:2507.17520.	927
873	language-action models. In <i>ICML</i> .		
874	Mohit Shridhar, Jesse Thomason, Daniel Gordon,	Cheng Yin, Yankai Lin, Wang Xu, Sikyuen Tam,	928
875	Yonatan Bisk, Winson Han, Roozbeh Mottaghi, Luke	Xiangrui Zeng, Zhiyuan Liu, and Zhouping Yin.	929
876	Zettlemoyer, and Dieter Fox. 2020. Alfred: A bench-	2025. Deepthinkvla: Enhancing reasoning capabil-	930
877	mark for interpreting grounded instructions for ev-	ity of vision-language-action models. <i>arXiv preprint</i>	931
878	eryday tasks. In <i>Proceedings of the IEEE/CVF con-</i>	<i>arXiv:2511.15669</i> .	932
879	<i>ference on computer vision and pattern recognition</i> ,		
880	pages 10740–10749.	Xiaohua Zhai, Basil Mustafa, Alexander Kolesnikov,	933
881	Mustafa Shukor, Dana Aubakirova, Francesco Capuano,	and Lucas Beyer. 2023. Sigmoid loss for language	934
882	Pepijn Kooijmans, Steven Palma, Adil Zouitine,	image pre-training. In <i>ICCV</i> , pages 11941–11952.	935
883	Michel Aractingi, Caroline Pascal, Martino Russi,	IEEE.	936
884	Andres Marafioti, and 1 others. 2025. Smolvla: A		
885	vision-language-action model for affordable and effi-	Huaixiu Steven Zheng, Swaroop Mishra, Xinyun Chen,	937
886	cient robotics. <i>arXiv preprint arXiv:2506.01844</i> .	Heng-Tze Cheng, Ed H. Chi, Quoc V. Le, and Denny	938
		Zhou. 2024. Take a step back: Evoking reasoning	939
887	Wenxuan Song, Jiayi Chen, Pengxiang Ding, Han Zhao,	via abstraction in large language models. In <i>ICLR</i> .	940
888	Wei Zhao, Zhide Zhong, Zongyuan Ge, Zhijun Li,		
889	Donglin Wang, Lujia Wang, Jun Ma, and Haoang Li.	Yifan Zhong, Fengshuo Bai, Shaofei Cai, Xuchuan	941
890	2025. PD-VLA: accelerating vision-language-action	Huang, Zhang Chen, Xiaowei Zhang, Yuanfei Wang,	942
891	model integrated with action chunking via parallel	Shaoyang Guo, Tianrui Guan, Ka Nam Lui, and 1	943
892	decoding. In <i>IROS</i> , pages 13162–13169. IEEE.	others. 2025. A survey on vision-language-action	944
		models: An action tokenization perspective. <i>arXiv</i>	945
893	Gerald Jay Sussman. 1975. <i>A computer model of skill</i>	<i>preprint arXiv:2507.01925</i> .	946
894	<i>acquisition</i> . Elsevier Science Inc.		
895	Hongyu Wang, Chuyan Xiong, Ruiping Wang, and Xilin	Brianna Zitkovich, Tianhe Yu, Sichun Xu, Peng Xu,	947
896	Chen. 2025. Bitvla: 1-bit vision-language-action	Ted Xiao, Fei Xia, Jialin Wu, Paul Wohlhart, Ste-	948
		fan Welker, Ayzaan Wahid, Quan Vuong, Vincent	949
		Vanhoecke, Huong T. Tran, Radu Soricut, Anikait	950
		Singh, Jaspiar Singh, Pierre Sermanet, Pannag R.	951

Sanketi, Grecia Salazar, and 35 others. 2023. RT-2: vision-language-action models transfer web knowledge to robotic control. In *CoRL*, volume 229 of *Proceedings of Machine Learning Research*, pages 2165–2183. PMLR.

A Technical Appendix

Appendix Table of Contents

1. Anticipatory defense on the choice of test environments and metrics (§ B)

Justification for using MINI-BEHAVIOR over alternatives (e.g., iGibson, CALVIN) based on controllability, symbolic abstraction, and isolation of language-grounding challenges.

2. Symbolic Representation of the Embodied Task (§ C) Formal definition of states, actions, and planning problems using classical STRIPS/PDDL semantics.

3. Details of the MINI-BEHAVIOR Environment (§ D) Overview of scene layout, object types, action space, and perception model.

4. Details of the MINI-BEHAVIOR-GRAN Extension (§ E) Procedure for generating multi-granularity instructions, details on the feature set, oracle instructor, and dataset construction.

5. Rule Set and Natural Language Prompt Examples (§ F) Details of the rule set bank of 20 tasks and sample generated instructions at different granularities.

6. Implementation Details (§ G) Model architecture (Prismatic-7B backbone, SigLIP+DINOv2 visual encoders), training hyperparameters, optimizer settings. VLM Predictor architecture and training details.

7. Additional Experimental Results (§ H) Extended analysis and ablations.

B Anticipatory defense on the choice of test environments and metrics

In this work, we selected Mini-BEHAVIOR as our testbed. This choice is motivated by several key considerations aligned with our research objectives. Our research objective is not to solve the full stack of robotic embodiment (which includes perception, continuous control, and physics), but specifically to isolate and analyze the impact of instruction granularity on language grounding and long-horizon planning.

Therefore, high-fidelity simulators (e.g., iGibson (Li et al., 2022a), BEHAVIOR-1K (Li et al., 2022b)) introduce significant noise through complex texture rendering and continuous motor control. While realistic, these factors act as confounders. When an agent fails in such environments, it is often difficult to attribute the failure to a lack of linguistic understanding (granularity issues) or a failure in low-level execution (e.g., grasping physics or occlusion). Mini-BEHAVIOR retains the high-level decision complexity of household tasks—long horizons, multi-object interactions, and state dependencies—while abstracting away low-level sensorimotor noise.

In the perspective of fast prototyping and iterative experimentation, we have to also consider the rendering and physics simulation costs. Mini-BEHAVIOR’s grid-based discrete control and simplified visual representation allow for rapid data generation and model training.

Another consideration is the diversity of the task structures so that we can systematically vary instruction granularity across a wide range of scenarios. As such, Crafter (Hafner, 2021), which focuses on survival in a single domain, lacks the diverse task structures needed to obtain robust insights about instruction granularity.

Table 3 provides a detailed comparison of Mini-BEHAVIOR against other prominent benchmarks, highlighting why they were less suitable for this specific investigation.

C Symbolic Representation of the Embodied Task

We formalize the Mini-BEHAVIOR environment as a deterministic grid-based symbolic world. This allows us to represent the challenge as a classical planning problem $\mathcal{P} = \langle \mathcal{F}, \mathcal{A}, s_0, G \rangle$, where:

- $\mathcal{F}$ is a finite set of propositional fluents representing the state variables of the environment (e.g., object locations, states like *cooked* or *open*).
- $\mathcal{A}$ is a finite set of actions available to the agent.
- $s_0 \subseteq \mathcal{F}$ is the initial state configuration.
- $G \subseteq \mathcal{F}$ is the goal condition that must be satisfied.

A state $s \in \mathcal{S}$ is defined as a truth valuation over $\mathcal{F}$, with the state space $\mathcal{S} \subseteq 2^{\mathcal{F}}$ encompassing all valid combinations of fluent truth values.

Table 3: Comparison of Embodied AI Environments regarding their suitability for studying Instruction Granularity.

Environment	Control	Horizon	Visuals	Symbolic Abstraction	Limitation
Mini-BEHAVIOR(*)	Discrete	Long	Grid	Support	Ideal balance: Supports symbolic state abstraction; diverse task suite; removes control noise while retaining logic complexity.
Crafter	Discrete	Long	Cartoon	Partial	Single task suite (survival); partially observable (memory heavy); few object types.
Meta-World	Continuous	Short	3D	No	Continuous control; no symbolic abstraction; focus on multi-task RL/manipulation, not long-horizon planning.
LIBERO	Continuous	Short	3D	No	Robotic arm focus; continuous control; no symbolic abstraction; VLA success rate already saturated (tasks relatively easy).
RoboTwin 2.0	Continuous	Short	3D	No	Continuous control; no symbolic abstraction; dual-arm manipulation focus; hard to segment expert trajectories and construct granular instructions.
CALVIN	Continuous	Medium	3D	No	Hard to design instruction granularity due to continuous action space (though provides diverse language); confounds planning with control difficulties.
ALFRED	Discrete	Long	3D Photo	Support	3D partially observable; poor infrastructure support for VLA training; high rendering cost.
iGibson / BEHAVIOR-1K	Continuous	Long	3D Photo	Partial	Good for long-horizon tasks, but 3D realistic vision context becomes a confounder.

D Details of MINI-BEHAVIOR Environment

To operationalize our formalism, we developed a **Universal PDDL Domain Model** (available at [minibehavior_gran_domain_model.pddl](#) in the codebase) capable of supporting all 20 household tasks without modification. This unified model comprises over 1,000 lines of PDDL 2.1 definitions and is fully compatible with the LAMA planner. We highly recommend checking the codebase for the complete domain model.

Table 4: Planning complexity in typical tasks.

Metric	Typical Range
Objects	15–40
Grid Locations	20–60
Ground Actions	10^3 – 10^5
Plan Length	50–200 primitive actions
Planning Time (LAMA)	0.5–300+ seconds

Search Space Characteristics

E Details of MINI-BEHAVIOR-GRAN Extension

MINI-BEHAVIOR provides the foundational task skeletons, which we extend in MINI-BEHAVIOR-GRAN to support multi-granularity instruction generation. A core component of this extension is the definition of a symbolic feature set used to generate rule-based policy sketches.

E.1 Feature Categories

The environment defines two primary categories of features: Boolean State Features and Quantitative Features.

E.1.1 Boolean State Features

These features evaluate whether a specific object instance satisfies a predicate. They require grounding (binding to specific object instances, e.g., `rag-01`).

E.1.2 Quantitative and Spatial Features

These features capture distances, counts, and spatial relationships. Crucially, they support temporal comparison allowing conditions such as “distance decreased” or “count dropped to zero.”

Table 5: Boolean State Features

Feature	Applicable Objects	Semantics
Holding	All Items	Is the agent holding the specified object?
isOpened	Cabinets, Drawers, Doors	Is the container/door open?
isToggled	Sinks, Stoves, Switches	Is the appliance turned on?
isSoaked	Rags, Sponges	Is the cleaning tool wet?
isCleaned	Shoes, Dishes, Floor	Is the object free of stains?
isDustFree	Furniture, Surfaces	Is the object free of dust?

Table 6: Available Count Functions for Rule Definitions

Category	Function Description
State-Based	count_not_cleaned: Objects that are currently stained. count_not_soaked: Cleaning tools that are dry. count_not_toggled: Appliances that are turned off. count_closed: Containers/doors that are closed. count_not_sliced: Food items that require slicing.
Spatial	count_not_on_floor: Objects not placed on the floor layer. count_isolated: Objects with no neighbors of the same type. count_not_inside_target: Objects not contained within a specific furniture. count_not_ontop: Objects not placed on top of a specific surface. count_not_near_target: Objects not within 1 step of a target type.
Task-Specific	count_incomplete_salad: Plates lacking necessary ingredients. count_not_complete_tables: Tables lacking specific setups (e.g., candles).

- **DistanceToNearest(type, condition):** Measures the Manhattan distance from the agent to the nearest object of a specific type.

– *Conditions:* $=0$ (adjacent), >0 , increase, decrease, change.

- **Count(type, condition, function):** Counts the number of objects of a specific type that satisfy a filter function.

– *Conditions:* $=0$, >0 , increase, decrease.

E.2 Count Functions

A diverse set of counting functions is implemented to support complex logical conditions (e.g., “all shoes are clean” is equivalent to *count_not_cleaned* $== 0$).

E.3 Reference Plan Generation and Task Complexity Analysis

While the original MINI-BEHAVIOR framework provides the high-level skeletons for household activities, it does not provide reference plans by default. Thus, we utilize the BFS (Best-First Width Search) classical planner to generate reference plans for all 20 tasks in MINI-BEHAVIOR-GRAN. These reference plans serve as the ground truth for our agents and allow us to quantitatively categorize the complexity of each task.

Detailed statistics regarding the specific rule sets used for instruction generation and the resulting natural language statistics are discussed in the next section.

F Rule Set and Natural Language Prompt Examples For Different Tasks

We present the width-0 rule set for 10 over 20 household tasks in the Mini-BEHAVIOR-Gran environment. For $w > 0$, please refer to the codebase at [mini-behavior-gran/mini_behavior_utils/policy_sketch/](https://github.com/mini-behavior-gran/mini_behavior_utils/policy_sketch/).

F.1 Task 1: Laying Wood Floors

Features:

- H : holding an isolated wood (plywood)
- p : distance to the nearest isolated wood
- t : distance to an arbitrary (non-held) wood we’ll place next to
- n : number of isolated woods

$\{\neg H, p > 0, n > 0\} \mapsto \{p \downarrow, t?\}$; move close to isolated wood
 $\{\neg H, p = 0, n > 0\} \mapsto \{H\}$; pick up isolated wood when reachable
 $\{H, t > 0, n > 0\} \mapsto \{t \downarrow\}$; move to some wood location
 $\{H, n > 0, t = 0\} \mapsto \{H?, n \downarrow, p?\}$; drop the isolated wood next to other wood

F.2 Task 2: Preparing Salad

Features:

- H_n : holding a not-sliceable food
- H_s : holding a sliceable food
- H_k : holding a slicer
- U : number of sliceable food that is not sliced yet
- d_p : distance to the selected incomplete plate p
- d_n, d_s : distances to nearest not-sliceable / sliceable
- k : distance to a slicer (knife)
- B : selected plate p complete bottom salad
- n : number of incomplete plates
- nn : number of plates that do not have bottom salad yet

$\{U > 0, \neg H_k, k > 0\} \mapsto \{k \downarrow\}$; move to slicer
 $\{U > 0, \neg H_k, k = 0\} \mapsto \{H_k\}$; hold a slicer
Acquire slicer: $\{U > 0, H_k, d_s > 0\} \mapsto \{d_s \downarrow\}$; move to sliceable food
 $\{U > 0, H_k, d_s = 0\} \mapsto \{U \downarrow, \neg H_k\}$; cut sliceable food
 $\{U = 0, H_k\} \mapsto \{\neg H_k\}$; put down slicer
 $\{n > 0, U = 0, nn > 0, \neg B, \neg H_n, d_n > 0\} \mapsto \{d_n \downarrow\}$
Place bottom salad: $\{n > 0, U = 0, nn > 0, \neg B, \neg H_n, d_p > 0\} \mapsto \{d_p \downarrow\}$
 $\{n > 0, U = 0, nn > 0, \neg B, H_n, d_p = 0\} \mapsto \{B, nn \downarrow, \neg H_n\}$
 $\{nn = 0, n > 0, U = 0, \neg H_s, d_s > 0\} \mapsto \{d_s \downarrow\}$; move to sliceable food
Handle sliceable salad: $\{nn = 0, n > 0, U = 0, \neg H_s, d_s = 0\} \mapsto \{H_s\}$; pick up sliceable food
 $\{nn = 0, n > 0, U = 0, H_s, d_p > 0\} \mapsto \{d_p \downarrow\}$; move to the plate
 $\{nn = 0, n > 0, U = 0, H_s, d_p = 0\} \mapsto \{n \downarrow, \neg H_s\}$; put down sliced food

F.3 Task 3: Cleaning Up the Kitchen Only

Features:

- *Holding flags:* H_b (blender), H_f (food), H_s (soap), H_r (rag), H_p (plate), H_o (vegetable oil)




Figure 10: Distribution of reference plan steps required for tasks in MINI-BEHAVIOR-GRAN obtained by BWFS classical planner.

- *Distances:* d_{ct} (countertop), d_{fr} (fridge), d_{si} (sink), d_{c1} (cabinet for plate), d_{c2} (cabinet for oil), d_x (to object x)
- *State booleans:* $\text{Open}(c)$, $\text{Tog}(s)$, Soaked , $\text{DustFree}(c)$, $\text{Clean}(p)$, $\text{OilIn}(c_2)$, Diff
- *Counters:* N_b (blenders not on countertop), N_a (apples not in fridge), N_{cas} (casseroles not in fridge), N_s (soaps not next to sink), N_{cab} (cabinets not dusted), N_r (rags not next to sink), N_{p_clean} (plates not clean), N_{oil} (oil not in cabinet), N_{p_store} (plates not stored), N_{close} (closed furniture)

$\{\neg H_r, d_{rag} > 0\} \mapsto \{d_{rag} \downarrow\}$; get rag
 $\{\neg H_r, d_{rag} = 0\} \mapsto \{H_r\}$
Cabinet dusting: $\{H_r, N_{cab} > 0, d_c > 0\} \mapsto \{d_c \downarrow\}$; clean cabinet
 $\{H_r, N_{cab} > 0, d_c = 0\} \mapsto \{N_{cab} \downarrow, \neg H_r\}$
 $\{N_{cab} = 0, N_{p_clean} = 0, \neg H_r, d_{rag} > 0\} \mapsto \{d_{rag} \downarrow\}$
 $\{N_{cab} = 0, N_{p_clean} = 0, \neg H_r, d_{rag} = 0\} \mapsto \{H_r\}$
Rag near/inside sink: $\{N_{cab} = 0, N_{p_clean} = 0, H_r, N_r > 0, d_{si} > 0\} \mapsto \{d_{si} \downarrow\}$
 $\{N_{cab} = 0, N_{p_clean} = 0, H_r, N_r > 0, d_{si} = 0\} \mapsto \{N_r \downarrow, \neg H_r\}$
 $\{N_{cab} = 0, N_{p_clean} = 0, \neg H_s, d_{soap} > 0\} \mapsto \{d_{soap} \downarrow\}$
 $\{N_{cab} = 0, N_{p_clean} = 0, \neg H_s, d_{soap} = 0\} \mapsto \{H_s\}$
Soap next to sink: $\{N_{cab} = 0, N_{p_clean} = 0, H_s, N_s > 0, d_{si} > 0\} \mapsto \{d_{si} \downarrow\}$
 $\{N_{cab} = 0, N_{p_clean} = 0, H_s, N_s > 0, d_{si} = 0\} \mapsto \{N_s \downarrow, \neg H_s\}$

F.4 Task 4: Organizing File Cabinet

Features:

- H_m : holding a marker
- H_f : holding a folder
- H_d : holding a document
- p_m : distance to the nearest marker
- p_f : distance to the nearest folder
- p_d : distance to the nearest document
- p_t : distance to a table
- p_c : distance to a valid cabinet target
- N_m : number of markers not on a table
- N_f : number of folders not in a cabinet
- N_d : number of documents not in a cabinet
- N_o : number of closed cabinets

Open Cabinet and Fridge: $\{N_{close} > 0, d_c > 0\} \mapsto \{d_c \downarrow\}$; move toward closed furniture
 $\{N_{close} = 0, d_c = 0\} \mapsto \{N_{close} \downarrow\}$
 $\{N_{close} = 0, \neg H_o, d_{oil} > 0\} \mapsto \{d_{oil} \downarrow\}$; approach oil
 $\{N_{close} = 0, \neg H_o, d_{oil} = 0\} \mapsto \{H_o\}$; pick up oil
Oil placement: $\{N_{close} = 0, H_o, N_{oil} > 0, d_{c2} > 0\} \mapsto \{d_{c2} \downarrow\}$; go to cabinet c_2
 $\{N_{close} = 0, H_o, N_{oil} > 0, d_{c2} = 0\} \mapsto \{N_{oil} \downarrow, \neg H_o\}$; place oil in c_2
 $\{\neg \text{Tog}(s), d_{si} > 0\} \mapsto \{d_{si} \downarrow\}$; toggle sink on
 $\{\neg \text{Tog}(s), d_{si} = 0\} \mapsto \{\text{Tog}(s)\}$
 $\{\neg H_r, d_{rag} > 0\} \mapsto \{d_{rag} \downarrow\}$; get rag
 $\{\neg H_r, d_{rag} = 0\} \mapsto \{H_r\}$
Plate cleaning: $\{H_r, \neg \text{Soaked}, d_{si} > 0\} \mapsto \{d_{si} \downarrow\}$; soak rag
 $\{H_r, \neg \text{Soaked}, d_{si} = 0\} \mapsto \{\text{Soaked}\}$
 $\{H_r, \text{Soaked}, N_{p_clean} > 0, d_{plate} > 0\} \mapsto \{d_{plate} \downarrow\}$
 $\{H_r, \text{Soaked}, N_{p_clean} > 0, d_{plate} = 0\} \mapsto \{N_{p_clean} \downarrow, d_{plate}?\}$
Plate placement distinct from oil: $\{N_{close} = 0, N_{oil} = 0, N_{p_clean} = 0, \neg H_p, N_{p_store} > 0, d_{plate} > 0\} \mapsto \{d_{plate} \downarrow\}$
 $\{N_{close} = 0, N_{oil} = 0, N_{p_clean} = 0, H_p, N_{p_store} > 0, d_{c1} > 0\} \mapsto \{d_{c1} \downarrow\}$
 $\{N_{close} = 0, N_{oil} = 0, N_{p_clean} = 0, H_p, N_{p_store} > 0, d_{c1} = 0\} \mapsto \{N_{p_store} \downarrow, \neg H_p\}$
 $\{N_b > 0, \neg H_b, d_{blender} > 0\} \mapsto \{d_{blender} \downarrow\}$
Blender to countertop: $\{N_b > 0, \neg H_b, d_{blender} = 0\} \mapsto \{H_b\}$
 $\{H_b, N_b > 0, d_{ct} > 0\} \mapsto \{d_{ct} \downarrow\}$
 $\{H_b, N_b > 0, d_{ct} = 0\} \mapsto \{N_b \downarrow, \neg H_b\}$
 $\{N_{close} = 0, \neg H_f, d_{apple} > 0\} \mapsto \{d_{apple} \downarrow\}$
Apple to fridge: $\{N_{close} = 0, \neg H_f, d_{apple} = 0\} \mapsto \{H_f\}$
 $\{N_{close} = 0, H_f, N_a > 0, d_{fr} > 0\} \mapsto \{d_{fr} \downarrow\}$
 $\{N_{close} = 0, H_f, N_a > 0, d_{fr} = 0\} \mapsto \{N_a \downarrow, \neg H_f\}$
 $\{N_{close} = 0, \neg H_f, d_{casserole} > 0\} \mapsto \{d_{casserole} \downarrow\}$
Casserole to fridge: $\{N_{close} = 0, \neg H_f, d_{casserole} = 0\} \mapsto \{H_f\}$
 $\{N_{close} = 0, H_f, N_{cas} > 0, d_{fr} > 0\} \mapsto \{d_{fr} \downarrow\}$
 $\{N_{close} = 0, H_f, N_{cas} > 0, d_{fr} = 0\} \mapsto \{N_{cas} \downarrow, \neg H_f\}$

Opening Cabinet: $\{N_o > 0, p_c > 0\} \mapsto \{p_c \downarrow\}$; move toward a closed cabinet
 $\{N_o > 0, p_c = 0\} \mapsto \{N_o \downarrow\}$; open the cabinet when reachable
 $\{N_m > 0, \neg H_m, p_m > 0\} \mapsto \{p_m \downarrow, p_t?\}$; move toward the nearest marker
Marker task: $\{N_m > 0, \neg H_m, p_m = 0\} \mapsto \{H_m, p_t?\}$; pick up the marker when reachable
 $\{N_m > 0, H_m, p_t > 0\} \mapsto \{p_t \downarrow\}$; move toward a table
 $\{N_m > 0, H_m, p_t = 0\} \mapsto \{N_m \downarrow, \neg H_m, p_m?\}$; put marker on table
 $\{N_f > 0, N_o = 0, \neg H_f, p_f > 0\} \mapsto \{p_f \downarrow, p_c?\}$
Folder task: $\{N_f > 0, N_o = 0, \neg H_f, p_f = 0\} \mapsto \{H_f, p_c?\}$
 $\{N_f > 0, N_o = 0, H_f, p_c > 0\} \mapsto \{p_c \downarrow\}$
 $\{N_f > 0, N_o = 0, H_f, p_c = 0\} \mapsto \{N_f \downarrow, \neg H_f, p_f?\}$

Task	Split	Class	Goal Description	Mean Steps
opening_packages	ID	short	Open every package.	10.2
installing_a_printer	ID	short	Place a printer on a table and turn it on.	10.6
laying_wood_floors	ID	short	Lay every plywood sheet on the floor, each touching at least one other sheet.	15.9
moving_boxes_to_storage	ID	short	Place every carton inside a shelf.	36.3
cleaning_a_car	ID	short	Make at least one car dust-free. Place a rag and soap together inside a bucket.	62.6
throwing_away_leftovers	ID	middle	Throw away every hamburger by placing it in a trash can.	68.3
organizing_file_cabinet	ID	middle	Put a marker on a table and file every folder and document in a cabinet.	75.5
making_tea	ID	middle	Slice a lemon. Put a teapot on an activated stove. Keep a soaked teabag with the teapot.	78.7
sorting_books	ID	middle	Place every book and every hardback on a shelf.	78.8
collect_misplaced_items	ID	middle	Place all gym shoes, necklaces, notebooks, and socks on a table.	90.7
thawing_frozen_food	ID	middle	Place a date label next to a fish. Put every fish next to a sink. Set an olive next to a sink.	101.7
boxing_books_up_for_storage	ID	long	Place every book inside a box.	134.4
storing_food	ID	long	Store every cookable item inside a cabinet.	144.6
preparing_salad	ID	long	Slice all apples and tomatoes; place them on plates. Put every radish and lettuce on plates. Ensure each plate holds at least one cookable item.	146.7
putting_away_dishes_after_cleaning	ID	long	Put every plate inside a cabinet.	147.7
cleaning_up_the_kitchen_only	ID	long	Put every blender on a countertop. Refrigerate all apples and casseroles. Keep soap next to a sink. Make plates unstained and cabinets dust-free. Place each rag either beside or inside the sink. Store plates and vegetable oil in different cabinets.	161.0
setting_up_candles	ID	long	On every table, place three distinct candles on top.	165.2
cleaning_shoes	OOD	short	Leave a towel on the floor and make sure every shoe is unstained and dust-free.	54.2
watering_houseplants	OOD	middle	Water every potted plant until it is soaked.	79.4
washing_pots_and_pans	OOD	long	Clean every teapot, kettle, and pan so they're unstained, then store each inside a cabinet.	149.5

Table 7: Classification of tasks in MINI-BEHAVIOR-GRAN based on the median number of plan steps, as determined by the BWFS planner. Tasks are categorized into short, middle, and long classes to reflect varying planning complexity based on plan length.

Document task: $\{N_d > 0, N_o = 0, \neg H_d, p_d > 0\} \mapsto \{p_d \downarrow, p_c?\}$
 $\{N_d > 0, N_o = 0, \neg H_d, p_d = 0\} \mapsto \{H_d, p_c?\}$
 $\{N_d > 0, N_o = 0, H_d, p_c > 0\} \mapsto \{p_c \downarrow\}$
 $\{N_d > 0, N_o = 0, H_d, p_c = 0\} \mapsto \{N_d \downarrow, \neg H_d, p_d?\}$

Olive Task: $\{\neg H_{olive}, p_{olive} > 0\} \mapsto \{p_{olive} \downarrow\}$
 $\{\neg H_{olive}, p_{olive} = 0\} \mapsto \{H_{olive}\}$
 $\{H_{olive}, p_{sink} > 0\} \mapsto \{p_{sink} \downarrow\}$
 $\{H_{olive}, p_{sink} = 0\} \mapsto \{N_{olive} \downarrow, \neg H_{olive}, p_{olive}?\}$
Date Task: $\{\neg H_{date}, N_{fish\ olive} = 0, p_{date} > 0\} \mapsto \{p_{date} \downarrow\}$
 $\{\neg H_{date}, N_{fish\ olive} = 0, p_{date} = 0\} \mapsto \{H_{date}\}$
 $\{H_{date}, N_{fish\ olive} = 0, p_{fish} > 0\} \mapsto \{p_{fish} \downarrow\}$
 $\{H_{date}, N_{fish\ olive} = 0, p_{fish} = 0\} \mapsto \{N_{date} \downarrow, \neg H_{date}, p_{date}?\}$

F.5 Task 5: Thawing Frozen Food

Features:

- H_{fish} : holding a fish
- H_{olive} : holding an olive
- H_{date} : holding a date
- p_{fish} : distance to the nearest fish
- p_{olive} : distance to the nearest olive
- p_{date} : distance to the nearest date
- p_{sink} : distance to a sink
- N_{fish} : number of fish not next to a sink
- N_{olive} : number of olives not next to a sink
- N_{date} : number of dates not next to a fish

Fish Task: $\{\neg H_{fish}, p_{fish} > 0\} \mapsto \{p_{fish} \downarrow\}$
 $\{\neg H_{fish}, p_{fish} = 0\} \mapsto \{H_{fish}\}$
 $\{H_{fish}, p_{sink} > 0\} \mapsto \{p_{sink} \downarrow\}$
 $\{H_{fish}, p_{sink} = 0\} \mapsto \{N_{fish} \downarrow, \neg H_{fish}, p_{fish}?\}$

F.6 Task 6: Making Tea

Features:

- H_l : holding a lemon
- H_{knife} : holding a knife
- H_t : holding a teapot
- H_{tb} : holding a tea bag
- p_l : distance to the nearest lemon
- p_{knife} : distance to the knife
- p_t : distance to the nearest teapot
- p_{tb} : distance to the nearest tea bag
- p_s : distance to the stove
- p_{sink} : distance to the sink
- is_sliced(lemon): whether the lemon is sliced
- is_toggled(stove): whether the stove is toggled on

- 1215 • `is_toggled(sink)`: whether the sink is toggled
- 1216 on
- 1217 • `is_soaked(tea_bag)`: whether the tea bag is
- 1218 soaked
- 1219 • `atsamelocation(tea_bag, teapot)`: tea bag at
- 1220 same location as teapot
- 1221 • `onTop(teapot, stove)`: teapot on stove

1222 *Toggle stove:* $\{\neg \text{is_toggled}(\text{stove}), p_s > 0\} \mapsto \{p_s \downarrow\}$
 $\{\neg \text{is_toggled}(\text{stove}), p_s = 0\} \mapsto \{\text{is_toggled}(\text{stove})\}$
 $\{\neg H_{knife}, p_{knife} > 0, \neg \text{is_sliced}(\text{lemon})\} \mapsto \{p_{knife} \downarrow\}$
 $\{\neg H_{knife}, p_{knife} = 0, \neg \text{is_sliced}(\text{lemon})\} \mapsto \{H_{knife}\}$
1223 *Slice lemon:* $\{H_{knife}, \neg \text{is_sliced}(\text{lemon}), p_l > 0\} \mapsto \{p_l \downarrow\}$
 $\{H_{knife}, \neg \text{is_sliced}(\text{lemon}), p_l = 0\} \mapsto \{\text{is_sliced}(\text{lemon})\}$
 $\{H_{knife}, \text{is_sliced}(\text{lemon})\} \mapsto \{\neg H_{knife}, p_{knife}?\}$
 $\{\neg \text{atsameloc}(\text{tb}, \text{tp}), \neg \text{onTop}(\text{tp}, \text{st}), \neg H_t, p_t > 0\} \mapsto \{p_t \downarrow\}$
 $\{\neg \text{atsameloc}(\text{tb}, \text{tp}), \neg \text{onTop}(\text{tp}, \text{st}), \neg H_t, p_t = 0\} \mapsto \{H_t\}$
 $\{\neg \text{atsameloc}(\text{tb}, \text{tp}), \neg \text{onTop}(\text{tp}, \text{st}), H_t, p_s > 0\} \mapsto \{p_s \downarrow\}$
1224 *Handle teapot:* $\{\neg \text{atsameloc}(\text{tb}, \text{tp}), \neg \text{onTop}(\text{tp}, \text{st}), H_t, p_s = 0\} \mapsto \{\text{onTop}(\text{tp}, \text{st}), \neg H_t, p_t?\}$
 $\{\neg \text{atsameloc}(\text{tb}, \text{tp}), \text{onTop}(\text{tp}, \text{st}), \neg H_{tb}, p_{tb} > 0\} \mapsto \{p_{tb} \downarrow\}$
 $\{\neg \text{atsameloc}(\text{tb}, \text{tp}), \text{onTop}(\text{tp}, \text{st}), \neg H_{tb}, p_{tb} = 0\} \mapsto \{H_{tb}\}$
 $\{\neg \text{atsameloc}(\text{tb}, \text{tp}), \text{onTop}(\text{tp}, \text{st}), H_{tb}, p_t > 0\} \mapsto \{p_t \downarrow\}$
 $\{H_{tb}, \neg \text{atsameloc}(\text{tb}, \text{tp}), \text{onTop}(\text{tp}, \text{st}), p_t = 0\} \mapsto \{\text{atsameloc}(\text{tb}, \text{tp}), \neg H_{tb}, p_{tb}?\}$

1225 F.7 Task 7: Opening Packages

1226 Features:

- 1227 • p : distance to the nearest package
- 1228 • N : number of unopened packages

1229 $\{N > 0, p > 0\} \mapsto \{p \downarrow\}$; move toward the nearest unopened package
 $\{N > 0, p = 0\} \mapsto \{N \downarrow\}$; open the package when reachable

1230 F.8 Task 8: Boxing Books Up for Storage

1231 Features:

- 1232 • H : holding an unboxed book
- 1233 • p : distance to the nearest unboxed book
- 1234 • b : distance to a valid box target
- 1235 • N : number of unboxed books

1236 $\{\neg H, p > 0\} \mapsto \{p \downarrow, b?\}$; move toward the nearest unboxed book
 $\{\neg H, p = 0\} \mapsto \{H\}$; pick up the book when reachable
 $\{H, N > 0, b > 0\} \mapsto \{b \downarrow\}$; move toward a chosen box spot
 $\{H, N > 0, b = 0\} \mapsto \{\neg H, N \downarrow, p?\}$; place the book into the box

1237 F.9 Task 9: Collect Misplaced Items

1238 Features:

- 1239 • H_s : holding a gym shoe
- 1240 • H_n : holding a necklace
- 1241 • H_b : holding a notebook
- 1242 • H_k : holding a sock
- 1243 • p_s : distance to the nearest gym shoe
- 1244 • p_n : distance to the nearest necklace
- 1245 • p_b : distance to the nearest notebook
- 1246 • p_k : distance to the nearest sock
- 1247 • p_t : distance to a table
- 1248 • N_s : number of gym shoes not on a table
- 1249 • N_n : number of necklaces not on a table
- 1250 • N_b : number of notebooks not on a table

- N_k : number of socks not on a table

$\{\neg H_s, N_s > 0, p_s > 0\} \mapsto \{p_s \downarrow, p_t?\}$
 $\{\neg H_s, N_s > 0, p_s = 0\} \mapsto \{H_s, p_t?\}$
Gym shoe task: $\{H_s, N_s > 0, p_t > 0\} \mapsto \{p_t \downarrow\}$
 $\{H_s, N_s > 0, p_t = 0\} \mapsto \{N_s \downarrow, \neg H_s, p_s?\}$
 $\{\neg H_n, N_n > 0, p_n > 0\} \mapsto \{p_n \downarrow, p_t?\}$
Necklace task: $\{\neg H_n, N_n > 0, p_n = 0\} \mapsto \{H_n, p_t?\}$
 $\{H_n, N_n > 0, p_t > 0\} \mapsto \{p_t \downarrow\}$
 $\{H_n, N_n > 0, p_t = 0\} \mapsto \{N_n \downarrow, \neg H_n, p_n?\}$
 $\{\neg H_b, N_b > 0, p_b > 0\} \mapsto \{p_b \downarrow, p_t?\}$
Notebook task: $\{\neg H_b, N_b > 0, p_b = 0\} \mapsto \{H_b, p_t?\}$
 $\{H_b, N_b > 0, p_t > 0\} \mapsto \{p_t \downarrow\}$
 $\{H_b, N_b > 0, p_t = 0\} \mapsto \{N_b \downarrow, \neg H_b, p_b?\}$
 $\{\neg H_k, N_k > 0, p_k > 0\} \mapsto \{p_k \downarrow, p_t?\}$
Sock task: $\{\neg H_k, N_k > 0, p_k = 0\} \mapsto \{H_k, p_t?\}$
 $\{H_k, N_k > 0, p_t > 0\} \mapsto \{p_t \downarrow\}$
 $\{H_k, N_k > 0, p_t = 0\} \mapsto \{N_k \downarrow, \neg H_k, p_k?\}$

F.10 Task 10: Putting Away Dishes After Cleaning

Features:

- H : holding a plate
- p : distance to the nearest plate
- b : distance to a valid cabinet target
- N : number of unboxed plates
- `is_opened`: whether the cabinet is opened

$\{\neg \text{is_opened}, b > 0\} \mapsto \{b \downarrow\}$; move toward the cabinet
 $\{\neg \text{is_opened}, b = 0\} \mapsto \{\text{is_opened}\}$; open the cabinet when reachable
 $\{\neg H, N > 0, p > 0, \text{is_opened}\} \mapsto \{p \downarrow, b?\}$
 $\{\neg H, N > 0, p = 0, \text{is_opened}\} \mapsto \{H, b?\}$
 $\{H, N > 0, b > 0, \text{is_opened}\} \mapsto \{b \downarrow\}$
 $\{H, N > 0, b = 0, \text{is_opened}\} \mapsto \{\neg H, N \downarrow, p?\}$

F.11 Instruction Generation Statistics

We used a template-based generator to translate the symbolic rule sets into natural language instructions, ensuring alignment with the VLA’s training data format. Table 8 presents the summary statistics for the generated instructions across the three granularity levels.

G Implementation Details

G.1 Model Architecture

Our policy is built upon the OpenVLA architecture (Kim et al., 2024). Specifically, we utilize the Prismatic-7B backbone (Karamcheti et al., 2024), which integrates a frozen Llama-2-7B language model with fused visual features. The visual encoder consists of a dual-stream setup combining SigLIP (for semantic understanding) and DINOv2 (for fine-grained spatial geometry) (Zhai et al., 2023; Oquab et al., 2024).

Table 8: Statistics of generated natural language instructions across granularity levels.

Granularity	Count	Avg Length (words)	Std Length	Min	Max	Unique Tokens	TTR (Type-Token Ratio)	Avg Sent. Length
Fine (F)	1636	41.12	11.36	19	87	140	0.0021	41.12
Medium (M)	859	38.62	12.76	22	81	108	0.0033	38.62
Coarse (C)	587	33.60	13.39	22	95	94	0.0048	33.60

G.2 Training Configuration

We fine-tune the model using Low-Rank Adaptation (LoRA) on the language model backbone, keeping the visual encoders frozen. All training setting variants (Pure, Centroid, Axial) share the same underlying training hyperparameters to ensure fair comparison.

Hardware and Duration: Training was conducted on a single node equipped with $8 \times$ NVIDIA A100 GPUs. A standard training run for 100,000 steps took approximately 23 hours to complete.

Table 9 summarizes the core hyperparameters used across our experiments.

Table 9: Hyperparameters for OpenVLA fine-tuning on MINI-BEHAVIOR-GRAN.

Hyperparameter	Value
Base Model	Prismatic-7B (Llama-2)
Visual Encoders	SigLIP + DINOv2
Parameter Efficient Tuning	LoRA (Rank $r = 32$)
Compute Hardware	$8 \times$ NVIDIA A100
Training Duration	~ 23 hours
Optimization Steps	100,000
Steps Before Decay	80,000
Learning Rate	2.5×10^{-4}
Batch Size (per GPU)	7 – 8

We also train a lightweight VLM predictor to quantify instruction predictability via perplexity. The model is fine-tuned using Low-Rank Adaptation (LoRA) on the Qwen3-VL-4B-Instruct backbone, with vision encoders frozen to preserve visual grounding capabilities established during pretraining. All models are trained on MINI-BEHAVIOR-GRAN instruction and image training data pairs, with training and evaluation splits maintained consistently across experiments. **Hardware and Duration:** Training was conducted on a single node equipped with $4 \times$ NVIDIA RTX 5090 GPUs. Each training run for 2 epochs. Table 10 summarizes the core hyperparameters used for VLM predictor training.

Table 10: Hyperparameters for Qwen3-VL-4B-Instruct fine-tuning on MINI-BEHAVIOR-GRAN (High-domain).

Hyperparameter	Value
Base Model	Qwen3-VL-4B-Instruct
Visual Encoders	Frozen (Qwen3-VL native)
Parameter Efficient Tuning	LoRA (Rank $r = 64$, $\alpha = 128$)
LoRA Target Modules	All linear layers
Compute Hardware	$4 \times$ NVIDIA RTX 5090
Training Epochs	2
Learning Rate	1.0×10^{-5}
Learning Rate Scheduler	Cosine with 10% warmup
Batch Size (per GPU)	32
Seed	[42,50,60]
Gradient Accumulation Steps	1
DeepSpeed Strategy	Zero-2

H Additional Experimental Results

This section provides additional detailed per-task performance breakdowns, and failure mode distributions.

H.1 Detailed Task Performance

Figure 11 details the success rates across all 20 tasks in the benchmark. We see some tasks are too difficult to solve and thus low across all instruction granularity. And then we also see the U shape trends in most tasks, showing the same phenomenon we discussed in (Q1) that both very fine and very coarse instructions are easier for VLA to follow.

H.2 Failure Analysis

Figure 12 presents the distribution of failure types across varying instruction widths. It confirms that low width encounters regression failures more frequently, while high width predominantly leads to stagnation failures. This aligns with our earlier observations in (Q3) regarding the distinct failure modes associated with different instruction granularities.




Figure 11: Detailed per-task success rates.




Figure 12: Failure type distribution across instruction widths.

H.3 More Plots for Width vs. Other Linguistic Heuristics as Metric for Language Grounding Complexity

As shown in the correlation heatmap Figure 13, w is the only metric that maintains a stable negative correlation with Success Rate across all horizons. In contrast, surface metrics like Token Count and Entity Count exhibit near-zero or even positive correlations in shorter horizons (e.g., $r = 0.09$ for 1-5 steps). This confirms the Complexity Paradox: while longer, more detailed instructions (higher token/entity counts) could actually provide helpful guidance that improves SR, increasing the latent planning width (w) invariably degrades performance.




Figure 13: Correlation heatmap between instruction width and other linguistic heuristics.

Similarly, the consistency analysis in Figure 14

shows that w possesses the lowest standard deviation ($\sigma = 0.033$), indicating its predictive reliability is invariant to task horizons. Surface metrics, particularly Action Verbs ($\sigma = 0.155$), are highly volatile. It also shows that w provides the highest mean predictive power ($|r| = 0.220$), more than doubling the predictive strength of Token Count (0.076) or Entity Count (0.072).




Figure 14: Consistency comparison between instruction width and other linguistic heuristics.

H.4 VLA Transformer Attention Weight Visualization




Figure 15: Attention from action tokens. Self-attention dominates, followed by vision; No clear trend across granularity levels.

Attention analysis (Figure 15) reveals there is no obvious reduction in language-token attention. This suggests that in VLA architectures, maintained attention weights do not necessarily equate to language grounding, highlighting a complex dynamic that requires future investigation.

Nevertheless, a closer inspection of the attention distributions reveals distinct architectural behaviors and compensatory mechanisms.

Autoregressive Models (Compensatory Reading). Autoregressive generation exhibits a heavy baseline reliance on historical action tokens (accounting for 0.74 of normalized attention in the Fine setting). Interestingly, when presented with Coarse instructions, attention to the dynamic instruction increases ($\uparrow 0.07$) while action attention decreases ($\downarrow 0.11$). This suggests a compensatory mechanism: as the rule width w increases and immediate subgoals disappear, the model attempts to ground more heavily on the language prompt to search for clues.

Parallel Decoding (Attending more on action modality). This architecture demonstrates a stark multimodal decoupling under coarse instructions. The attention paid to visual tokens drops significantly ($\downarrow 0.11$) in the Coarse setting, accompanied by an increased reliance on generated action tokens ($\uparrow 0.09$). This could imply that without detailed instruction guidance at each step, parallel decoding intend to trust more on action history to decide the next actions, showing a different strategy from autoregressive models.

Discrete Diffusion (Structural Stability). Diffusion models exhibit the most balanced attention distribution, maintaining the highest baseline attention to both dynamic instructions (0.20) and visual features (0.32). Notably, this distribution remains remarkably rigid regardless of instruction granularity (with fluctuations bounded within ± 0.04).

H.5 Sample Generated Instructions by Granularity

To illustrate the qualitative differences between granularity levels, we provide five random samples for each category below. Note how Low (L) granularity explicitly specifies navigation and preconditions (e.g., “move toward,” “not yet at”), while High (H) granularity focuses almost exclusively on the final state change (e.g., “reduce the count of...”), shown in Table 11.

Fine Granularity (F) Samples

1. At step 11: When the count of uncleaned plate is above 0, not yet at sink 0 and sink is off, please try to move toward sink 0.
2. At step 8: When not yet at soap 0, not holding soap 0, the count of unwiped car is 0 and the count of rag not inside bucket is 0, please try to move toward soap 0.
3. At step 27: When the count of chip and oatmeal sugar and vegetable oil not inside cabinet is above 0, you are at oatmeal 1 and not holding oatmeal 1, please pick up oatmeal 1.
4. At step 15: When the count of plate not inside cabinet is above 0, you are at plate 3, the count of closed cabinet is 0 and not holding plate 3, please pick up plate 3.
5. At step 13: When the count of fish and olive not near sink is above 0, not yet at olive 0 and not holding olive 0, please try to move toward olive 0.

Medium Granularity (M) Samples

1. At step 2: When the count of plate not inside cabinet is above 0, the count of closed cabinet is 0 and not holding plate 2, please pick up plate 2.
2. At step 6: When the count of uncleaned plate is above 0, the count of dry rag is above 0 and holding rag 0, please reduce the count of dry rag.
3. At step 9: When the count of plate not inside cabinet 0 is above 0, the count of closed cabinet is 0, the count of vegetable oil not inside cabinet 1 is 0, the count of uncleaned plate is 0 and not holding plate 0, please pick up plate 0.
4. At step 13: When the count of book not inside box is above 0 and not holding book 2, please pick up book 2.
5. At step 9: When the count of chip and oatmeal sugar and vegetable oil not inside cabinet is above 0, holding chip 1 and the count of closed cabinet is 0, please not holding chip 1, then reduce the count of chip and oatmeal sugar and vegetable oil not inside cabinet.

Coarse Granularity (C) Samples

1. At step 2: When the count of hamburger not inside ashcan is above 0, please reduce the count of hamburger not inside ashcan.
2. At step 5: When the count of plate not inside cabinet is above 0 and the count of closed cabinet is 0, please reduce the count of plate not inside cabinet.
3. At step 3: When the count of teapot not on stove is above 0, please reduce the count of teapot not on stove.
4. At step 9: When the count of plate not inside cabinet is above 0 and the count of closed cabinet is 0, please reduce the count of plate not inside cabinet.
5. At step 8: When the count of chip and oatmeal sugar and vegetable oil not inside cabinet is above 0 and the count of closed cabinet is 0, please reduce the count of chip and oatmeal sugar and vegetable oil not inside cabinet.

Table 11: Sample generated instructions across granularity levels.